

# SoilGrids250 — Import tiles via WCS

D G Rossiter

[david.rossiter@isric.org](mailto:david.rossiter@isric.org)

12-January-2025

## Table of Contents

Introduction .....	1
Directories.....	2
Packages.....	3
Coördinate Reference System (CRS).....	3
Parameters.....	3
Area of Interest (AOI).....	4
Destination directories.....	5
Accessing SoilGrids250 WCS with GDAL .....	6
Resample to WGS84 geographic coördinates.....	10
Save tile .....	12

## Introduction

**SoilGrids** is a system for global digital soil mapping that makes use of global soil profile information and covariate data to model the spatial distribution of soil properties across the globe. SoilGrids250 is a collections of soil property maps at six standard depths at 250 m grid resolution.

SoilGrids250 filenames, procedures etc. are explained in a [FAQ](#). The choice of the [Goode Homolosine projection](#) is explained in Moreira de Sousa, L., L. Poggio, and B. Kempen. 2019. *Comparison of FOSS4G Supported Equal-Area Projections Using Discrete Distortion Indicatrices. ISPRS International Journal of Geo-Information* 8(8): 351. <https://doi.org/10.3390/ijgi8080351>.

Download of SoilGrids layers is most convenient by Web Coverage Service (WCS), see [this description of procedures from ISRIC](#).

This script creates a tile for a property and depth slice, over a Area of Interest delimited by geographic coördinates. Tiles are created in EPSG 4326 (WGS84 long/lat) with SoilGrid's nominal 250 m grid resolution.

This tile can then be compared with other PSM products or combined into a raster stack for pattern analysis.

To use this script:

1. Adjust the [directory structure](#) to your system.

Steps 2–3 refer to the YAML headers of the R Markdown source document, or external calls with `knitr::render`. You can ask to be prompted for the parameters with the R command at the console:

```
rmarkdown::render("SoilGrids250_WCS_import.Rmd", output_format =  
"html_document",  
                params = "ask")
```

or you can specify them directly and run this script from the R console, e.g.,

```
rmarkdown::render("SoilGrids250_WCS_import.Rmd", output_format = NULL,  
                params = list(lrc_long = -76, lrc_lat = 42,  
                             size = 1, quantile.n = 4,  
                             voi.n = 3, depth.n = 1))
```

Here we specify `output_format = NULL` to just get the tile but not produce a document.

See the respective sections of the code for the parameter definitions.

2. [Select a property and quantile and a depth slice](#), using the YAML header or by knitting with parameters.
3. [Select an Area of Interest](#), typically a  $1 \times 1^\circ$  tile, using the YAML header or by knitting with parameters.
4. Either compile to HTML or PDF (“knit”), or “Run All” within R Markdown, or call `rmarkdown::render` from the R command line, as explained above.
5. The processed tile will be in the directory structure, in a [subdirectory named for the AOI](#).

## Directories

Set base directories, specific to the local file system.

1. `base.dir.import`: This is where downloaded large GeoTIFF are located. Because of their size they may be on a separate file system, e.g., removable or networked drive.
2. `base.dir`: This is where the processed SoilGrids250 maps are stored.

```
(base.dir.sg <- path.expand("~/ds_reference/DSM2025/"))
```

```
[1] "/Users/rossiter/ds_reference/DSM2025/"
```

```
(base.dir.sg.import <- path.expand("~/tmp/DSM2025/"))
[1] "/Users/rossiter/tmp/DSM2025/"
```

These are the base of destination directories built [below](#)

## Packages

```
library(sf)           # spatial data types
library(terra)        # raster data, replaces `raster`
require(rgdal)        # directly access GDAL functions -- deprecated but
still works
require(gdalUtilities)
require(XML)
```

GDAL is used for spatial data import/export, coördinate systems etc. Check the GDAL installation, and whether the WCS service is available.

```
gdal() # version
gdal.drivers <- terra::gdal(drivers = TRUE)
ix <- which(gdal.drivers$name == "WCS")
gdal.drivers[ix, ]
```

## Coördinate Reference System (CRS)

We want to use geographic coördinates for the tile. But the ISRIC WCS does not seem to serve this – or at least, I can not figure out how. So, we must download SoilGrids250 in the native Homolosine CRS. For this we need to know the bounding box in that CRS.

Note: This CRS with pseudo-EPSG code 152160 should be added to the epsg file of the PROJ database on your system <sup>1</sup> as a final line, as explained [here](#). But for now do not do this, just specify the projection directly.

```
crs.igh <- '+proj=igh +lat_0=0 +lon_0=0 +datum=WGS84 +units=m +no_defs'
```

## Parameters

Define the variables for the soil property and layer of interest. See [here](#) for the naming conventions

Q0.05 - 5% quantile from the Quantile Random Forest (QRF); Q0.5 - median of the distribution from the QRF – note *not* Q0.50; Q0.95 - 95% quantile from the QRF; mean - mean of the distribution.

```
quantile.list <- c("Q0.05", "Q0.5", "Q0.95", "mean")
```

---

<sup>1</sup> for example /Library/Frameworks/PROJ.framework/Versions/6/Resources/proj/epsg

Here are the properties predicted by SoilGrids250:

```
voi.list.sg <- c("clay", "silt", "sand", "phh2o", "cec", "soc", "bdod",  
"cfvo", "nitrogen", "ocd")
```

We do not include ocs (Organic C stocks) here because that is for the whole profile, not by depth slice.

Here are the depth slices predicted by SoilGrids250:

```
depth.list <- paste0(c("0-5", "5-15", "15-30", "30-60", "60-100", "100-  
200"), "cm")
```

Parameters for this run:

```
print(paste("VOI:", voi.list.sg[params$voi.n]))  
[1] "VOI: cfvo"  
print(paste("Depth slice:", depth.list[params$depth.n]))  
[1] "Depth slice: 0-5cm"  
print(paste("Quantile:", quantile.list[params$quantile.n]))  
[1] "Quantile: mean"
```

Set the property, depth and quantile from the YAML or rendering parameters:

```
voi <- voi.list.sg[params$voi.n]  
depth <- depth.list[params$depth.n]  
quantile.sg <- quantile.list[params$quantile.n]  
(voi_layer <- paste(voi, depth, quantile.sg, sep="_"))  
[1] "cfvo_0-5cm_mean"
```

## Area of Interest (AOI)

The AOI is a square tile using WGS84 geographic coordinates. A  $1 \times 1^\circ$  is typical.

Specify the *lower-right corner* and *tile size* from the YAML or rendering parameters:

```
tile.lrc <- c(params$lrc_long, params$lrc_lat) # Lower-right corner  
size <- params$size # tile dimensions  
print(paste("Lower-right corner: longitude", params$lrc_long,  
"; latitude", params$lrc_lat,  
"; tile size:", params$size, "degree square"))  
[1] "Lower-right corner: longitude 80 ; latitude 21 ; tile size: 1 degree  
square"
```

Compute the four corner and the bounding box. Note because of the projection this is somewhat larger than a  $1 \times 1^\circ$  tile.

```

tile.ulc <- c(tile.lrc[1]-size, tile.lrc[2]+size) # upper-left corner
m <- matrix(c(tile.lrc[1]-size, tile.lrc[2]+size, #ulc
              tile.lrc[1], tile.lrc[2]+size, #urc
              tile.lrc[1], tile.lrc[2], #lrc
              tile.lrc[1]-size, tile.lrc[2] #lcc
            ),
            nrow=4, byrow = TRUE)
m <- rbind(m, m[1,]) # close the polygon
bb.ll <- st_sfc(st_polygon(list(m)))

```

The CRS of this bounding box is WGS84 geographic coordinates (EPSG code 4326).

```

st_crs(bb.ll) <- 4326
# print(bb.ll)
# st_boundary(bb.ll)

```

A prefix for directories, to keep AOI results separate.

```

(AOI.dir.prefix <- paste0("lat", tile.lrc[2], tile.ulc[2],
                          "_lon", tile.ulc[1], tile.lrc[1]))

[1] "lat2122_lon7980"

```

## Destination directories

Set import and destination directories, adding to the base directories the variable of interest, quantile, depth. Make sure the directory exists.

```

(dest.dir.sg.import <- paste0(base.dir.sg.import, AOI.dir.prefix))

[1] "/Users/rossiter/tmp/DSM2025/lat2122_lon7980"

if (!dir.exists(dest.dir.sg.import)) {
  dir.create(dest.dir.sg.import, recursive = TRUE)
}
(dest.dir.sg <- paste0(base.dir.sg,
                      AOI.dir.prefix))

[1] "/Users/rossiter/ds_reference/DSM2025/lat2122_lon7980"

if (!dir.exists(dest.dir.sg)) {
  dir.create(dest.dir.sg, recursive = TRUE)
}

```

Convert the long/lat bounding box to the SoilGrids250 projection. We want the extreme values in both X and Y, to ensure we cover the whole tile. If we just use the corners we will cut off some parts at the upper-right and lower-left.

```

(bb.igh <- st_transform(bb.ll, crs.igh)) # reproject the polygon

```

```

Geometry set for 1 feature
Geometry type: POLYGON
Dimension:      XY
Bounding box:   xmin: 8397053 ymin: 2337709 xmax: 8535870 ymax: 2449029
Projected CRS: +proj=igh +lat_0=0 +lon_0=0 +datum=WGS84 +units=m +no_defs

(bb.igh.coords <- st_coordinates(bb.igh)[,1:2]) # convert to coördinates, we
only need 2D

      X      Y
[1,] 8397053 2449029
[2,] 8500266 2449029
[3,] 8535870 2337709
[4,] 8431944 2337709
[5,] 8397053 2449029

# convert to a bounding box, must order these as c(ulx, uly, lrx, lry)
bb.sg <- as.vector(c(min(bb.igh.coords[, "X"]),
                      max(bb.igh.coords[, "Y"]),
                      max(bb.igh.coords[, "X"]),
                      min(bb.igh.coords[, "Y"])))

```

## Accessing SoilGrids250 WCS with GDAL

Adapted from instructions by Luis de Souza (ISRIC), see [https://git.wur.nl/isric/soilgrids/soilgrids.notebooks/-/blob/master/markdown/wcs\\_from\\_R.md](https://git.wur.nl/isric/soilgrids/soilgrids.notebooks/-/blob/master/markdown/wcs_from_R.md).

The idea is to set up an XML file with the request and then feed that to GDAL commands.

Note on the XML file name: re-using the same file name if this is called from another script seems to somehow not refresh a server cache – unclear how this happens – but anyway, this is why we make each XML file with a temporary name.

```

wcs_request <- "DescribeCoverage"
wcs_path <- paste0("https://maps.isric.org/mapserv?map=/map/", voi, ".map") #
Path to the WCS. See maps.isric.org
wcs_service = "SERVICE=WCS"
wcs_version = "VERSION=2.0.1"
(wcs <- paste(wcs_path, wcs_service, wcs_version, wcs_request, sep="&"))

[1]
"https://maps.isric.org/mapserv?map=/map/cfvo.map&SERVICE=WCS&VERSION=2.0.1&DescribeCoverage"

l1 <- newXMLNode("WCS_GDAL")
l1.s <- newXMLNode("ServiceURL", wcs, parent=l1)
l1.l <- newXMLNode("CoverageName", voi_layer, parent=l1)
# Save to local disk
xml.out <- paste0(".", tempfile("sg", "", ".xml"))

```

```

tmp <- saveXML(l1, file = xml.out)
gdalinfo(xml.out)

Driver: WCS/OGC Web Coverage Service
Files: ./sg165a7cf630b7.xml
Size is 159246, 58034
Coordinate System is:
PROJCRS["Interrupted_Goode_Homolosine",
  BASEGEOGCRS["WGS 84",
    DATUM["World Geodetic System 1984",
      ELLIPSOID["WGS 84",6378137,298.257223563,
        LENGTHUNIT["metre",1]],
      ID["EPSG",6326]],
    PRIMEM["Greenwich",0,
      ANGLEUNIT["Degree",0.0174532925199433]]],
  CONVERSION["unnamed",
    METHOD["Interrupted Goode Homolosine"]],
  CS[Cartesian,2],
    AXIS["(E)",east,
      ORDER[1],
      LENGTHUNIT["metre",1,
        ID["EPSG",9001]]],
    AXIS["(N)",north,
      ORDER[2],
      LENGTHUNIT["metre",1,
        ID["EPSG",9001]]]]
Data axis to CRS axis mapping: 1,2
Origin = (-19949875.000000000000000,8361125.000000000000000)
Pixel Size = (250.000000000000000,-250.000000000000000)
Corner Coordinates:
Upper Left  (-19949875.000, 8361125.000)
Lower Left  (-19949875.000,-6147375.000)
Upper Right (19861625.000, 8361125.000)
Lower Right (19861625.000,-6147375.000)
Center      ( -44125.000, 1106875.000) ( 0d51'35.71"W, 9d56'35.62"N)
Band 1 Block=1024x512 Type=Int16, ColorInterp=Undefined
  NoData Value=-32768
  Overviews: 79623x29017, 39811x14508, 19905x7254, 9952x3627, 4976x1813,
2488x906, 1244x453, 622x226

file.remove(xml.out)

[1] TRUE

# also remove 'helper' XML if created by saveXML
xml.DC <- paste0(".", strsplit(xml.out,".", fixed = TRUE)[[1]][2], ".DC.xml")
if (file.exists(xml.DC)) file.remove(xml.DC)

[1] TRUE

```

Ignore the GDAL warnings, because we know the Homolosine projection used in SG250 is not in the GDAL database.

Get the coverage:

```
(wcs <- paste(wcs_path, wcs_service, wcs_version, sep="&"))

[1]
"https://maps.isric.org/mapserv?map=/map/cfvo.map&SERVICE=WCS&VERSION=2.0.1"

l1 <- newXMLNode("WCS_GDAL")
l1.s <- newXMLNode("ServiceURL", wcs, parent=l1)
l1.l <- newXMLNode("CoverageName", voi_layer, parent=l1)
xml.out <- paste0(" ", tempfile("sg", "", ".xml"))
tmp <- saveXML(l1, file = xml.out)
(file.out <- paste0(dest.dir.sg.import, "/", voi_layer, ".tif"))

[1] "/Users/rossiter/tmp/DSM2025/lat2122_lon7980/cfvo_0-5cm_mean.tif"

gdal_translate(xml.out, file.out,
  tr=c(250,250),
  projwin = bb.sg, projwin_srs = crs.igh, # corners in this CRS
  co=c("TILED=YES", "COMPRESS=DEFLATE", "PREDICTOR=2", "BIGTIFF=YES")
)
file.remove(xml.out)

[1] TRUE

# also remove 'helper' XML if created by saveXML
xml.DC <- paste0(" ", strsplit(xml.out, ".", fixed = TRUE)[[1]][2], ".DC.xml")
if (file.exists(xml.DC)) file.remove(xml.DC)

[1] TRUE
```

Again, ignore the GDAL warnings.

Check the result by reading into R, summarizing, and plotting. Note this is in the Homolosine projection used by SG250.

```
r.sg <- terra::rast(file.out)
print(r.sg)


class      : SpatRaster
dimensions : 445, 555, 1  (nrow, ncol, nlyr)
resolution : 250, 250  (x, y)
extent     : 8396875, 8535625, 2337875, 2449125  (xmin, xmax, ymin, ymax)
coord. ref.: Interrupted_Goode_Homolosine
source     : cfvo_0-5cm_mean.tif
name       : cfvo_0-5cm_mean

summary(values(r.sg))
```



```
cfvo_0-5cm_mean
Min.   : 40.0
1st Qu.:117.0
Median :149.0
Mean   :154.6
3rd Qu.:194.0
Max.   :314.0
NA's   :6980
```

```
terra::plot(r.sg, col=(sp::bpy.colors(50)), main = voi_layer)
```



### *SoilGrids IGH tile*

The 0 values in SoilGrids are in fact un-mapped areas, so convert these to NA.

```
quantile(values(r.sg), seq(0,1,by=0.05), na.rm = TRUE)
```

	0%	5%	10%	15%	20%	25%	30%	35%	40%	45%	50%	55%	60%	65%	70%
75%															
	40	80	89	99	109	117	125	131	137	143	149	154	161	169	181
194															
	80%	85%	90%	95%	100%										
	205	215	226	240	314										

```
values(r.sg) <- ifelse(values(r.sg) < 1, NA, values(r.sg))
quantile(values(r.sg), seq(0,1,by=0.05), na.rm = TRUE)
```

	0%	5%	10%	15%	20%	25%	30%	35%	40%	45%	50%	55%	60%	65%	70%
75%															
	40	80	89	99	109	117	125	131	137	143	149	154	161	169	181
194															
	80%	85%	90%	95%	100%										
	205	215	226	240	314										

```
# terra::plot(r.sg, col=(sp::bpy.colors(50)), main = voi_layer)
```

## Resample to WGS84 geographic coördinates

Set the CRS and resample to WGS84 geographic coördinates. Use bilinear interpolation – notice that the values are now fractional, rather than integers as in the original SoilGrids layer.

```
st_bbox(r.sg)
```

	xmin	ymin	xmax	ymax
	8396875	2337875	8535625	2449125

```
st_crs(r.sg)$proj4string
```

```
[1] "+proj=igh +lon_0=0 +x_0=0 +y_0=0 +datum=WGS84 +units=m +no_defs"
```

```
st_crs(r.sg)$epsg      # not in the EPSG database
```

```
[1] NA
```

```
r.sg.84 <- terra::project(r.sg, "epsg:4326")
```

```
quantile(values(r.sg.84), seq(0,1,by=.1), na.rm = TRUE)
```

	0%	10%	20%	30%	40%	50%	60%
70%							
	42.51252	90.41739	109.66060	124.78061	137.85013	148.78765	161.05038
	181.04806						
	80%	90%	100%				
	204.83351	224.92386	306.69016				

```
st_crs(r.sg.84)$proj4string
```

```
[1] "+proj=longlat +datum=WGS84 +no_defs"
```

```
st_crs(r.sg.84)$epsg
```

```
[1] 4326
```


Crop this to the long/lat bounding box – recall, the Homolosine bounding box is larger:

```
r.sg.84.crop <- crop(r.sg.84, bb.ll)
```

```
quantile(values(r.sg.84.crop), seq(0,1,by=0.05), na.rm = TRUE)
```

	0%	5%	10%	15%	20%	25%	30%
35%							
	42.51252	80.44707	89.43628	98.25671	107.97195	116.25165	123.78732
130.95309							
	40%	45%	50%	55%	60%	65%	70%
75%							
	137.24946	142.90861	148.33936	153.96141	160.28127	168.58617	180.75125
194.81840							
	80%	85%	90%	95%	100%		
	206.03329	215.62958	225.40694	237.50547	306.69016		

```
terra::plot(r.sg.84.crop, col=(sp::bpy.colors(50)), main = voi_layer)
```



*SoilGrids WGS84 tile*

## Save tile

Save this map for further processing, e.g., comparing with other DSM products.

Save the tile. Note that the file name includes the property name and depth slice. Specify the float-4 bit datatype and a GeoTIFF “world” file. Each tile is about 750 kB.

```
print(paste("file name:",  
           file.out <- paste0(dest.dir.sg, "/", voi_layer, '.tif')))
```

```
[1] "file name: /Users/rossiter/ds_reference/DSM2025/lat2122_lon7980/cfvo_0-  
5cm_mean.tif"
```

```
writeRaster(r.sg.84.crop, file = file.out,  
            overwrite=TRUE, datatype="FLT4S",  
            gdal = "TFW=YES")
```

This can now be imported into analysis scripts.